

mojaloop

# Performance - Coil

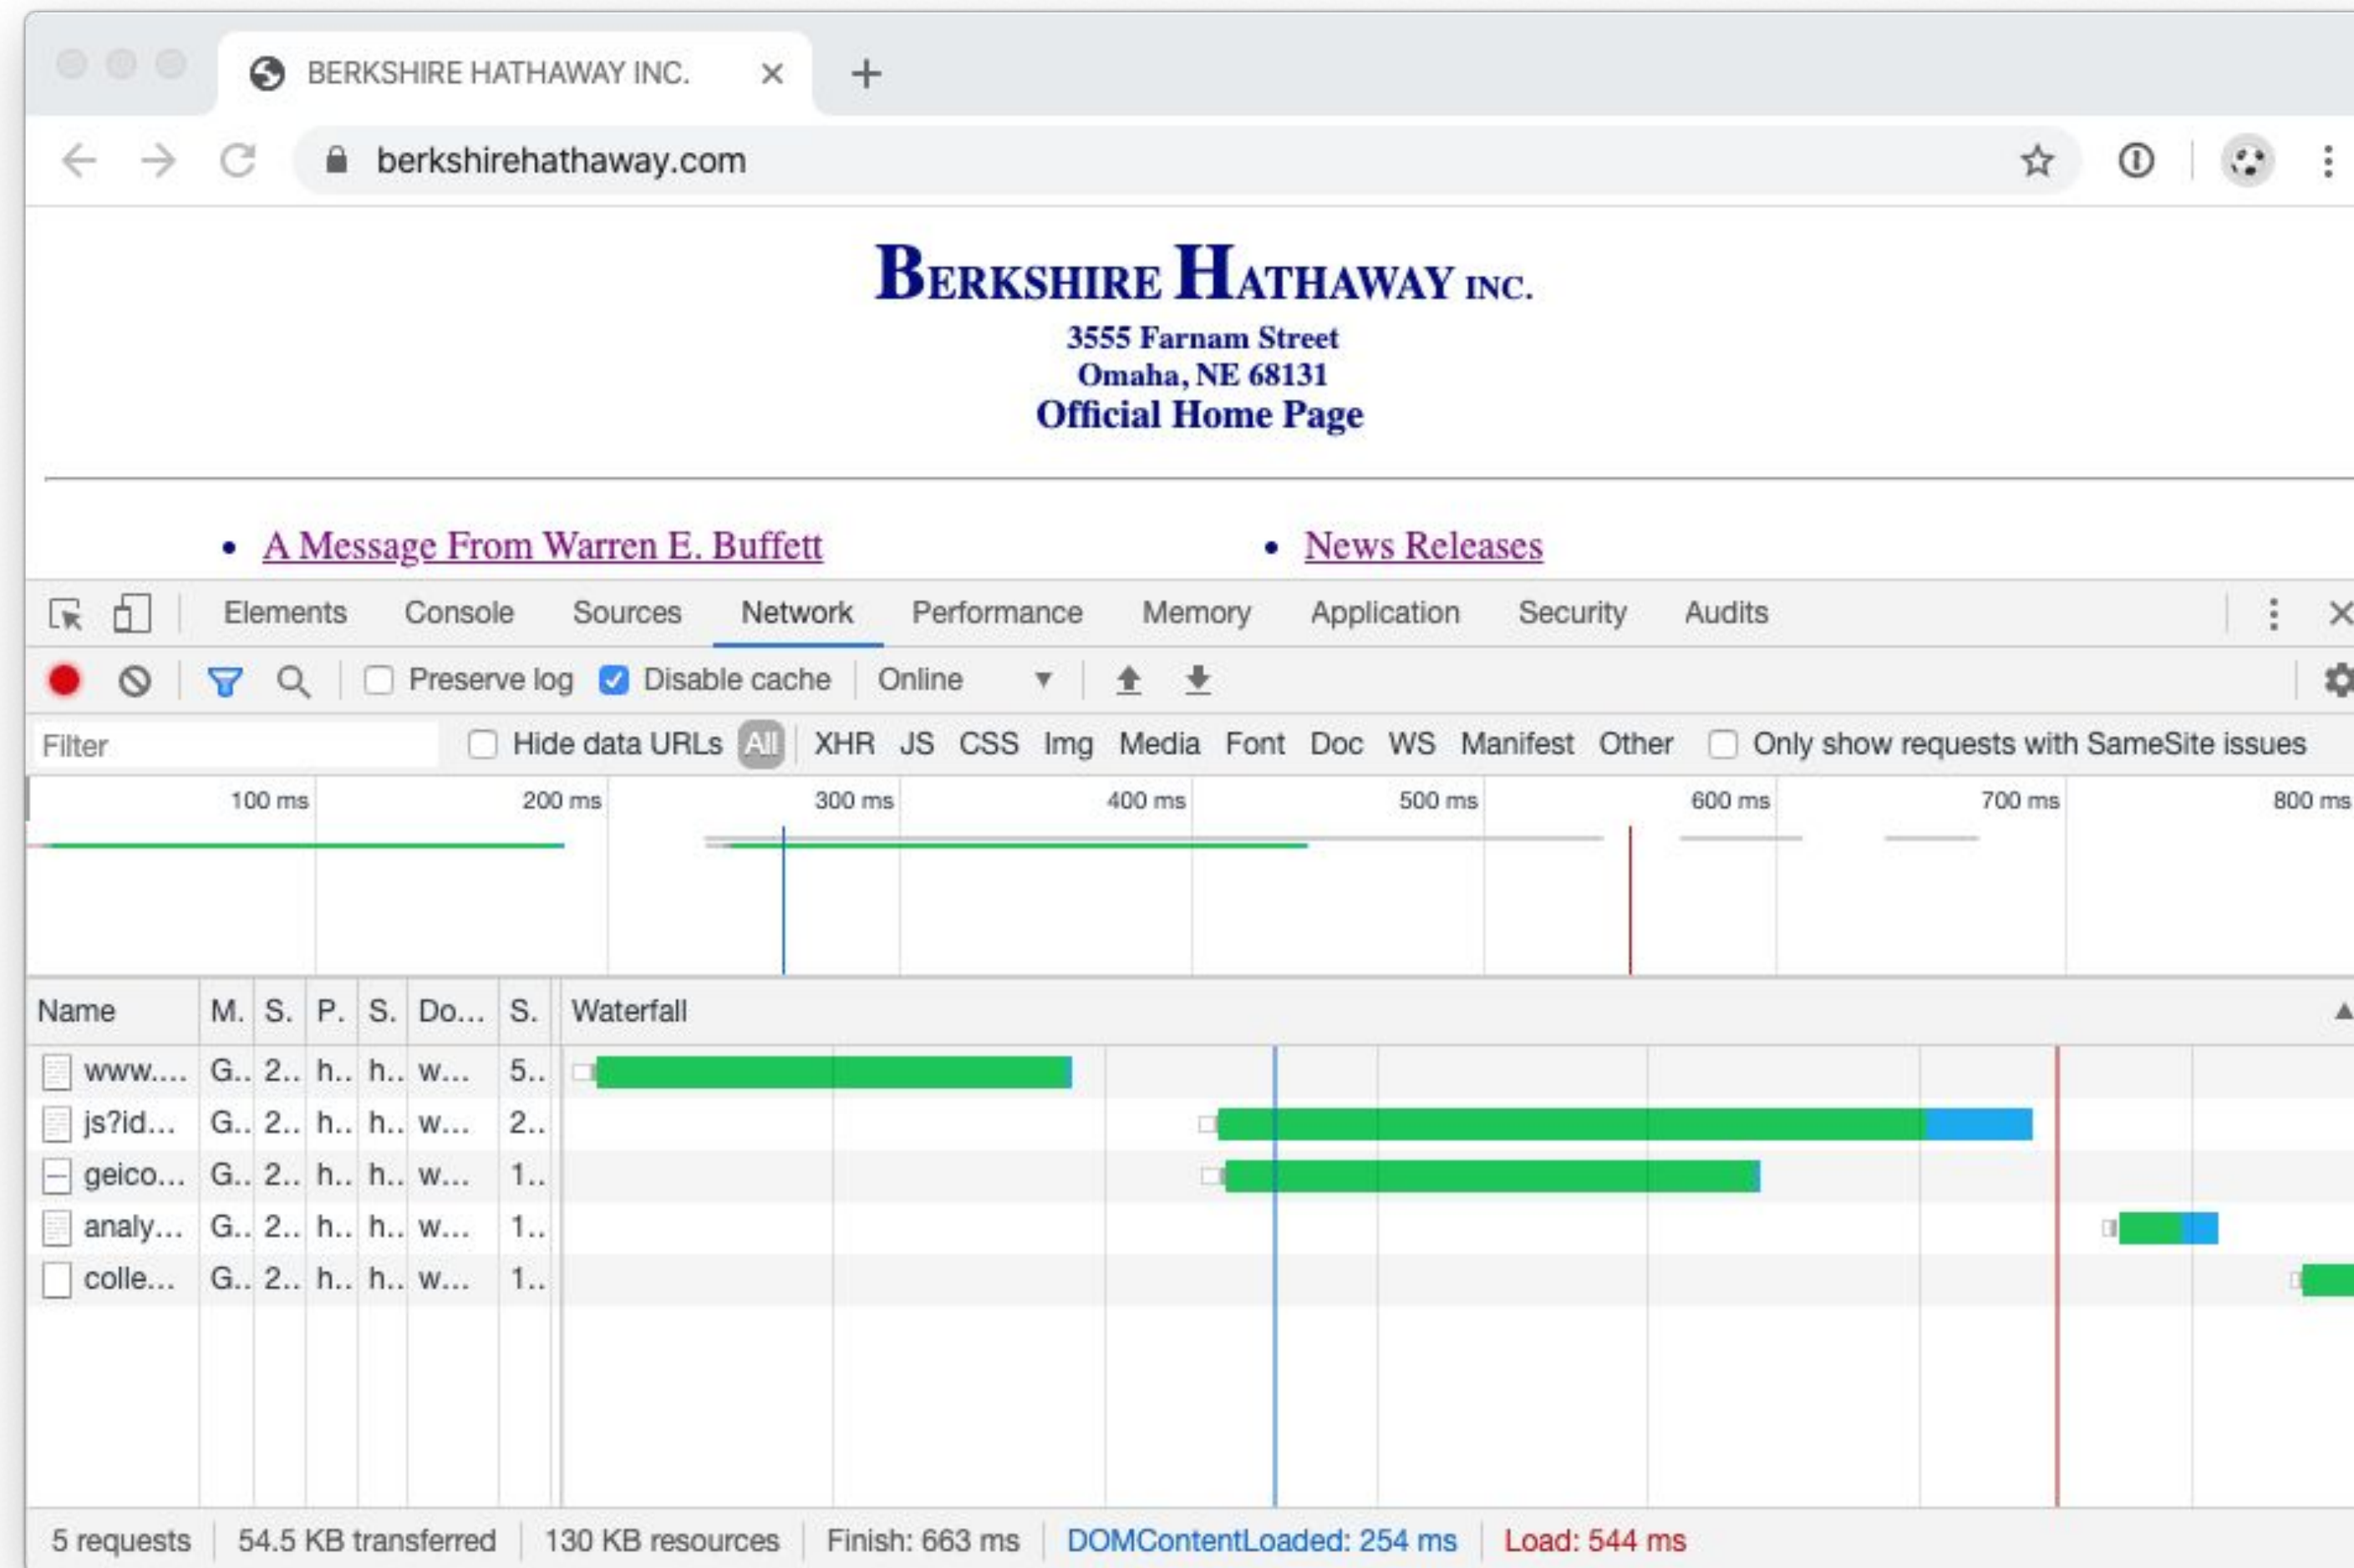
Time gaps and waterfalls

mojaloop

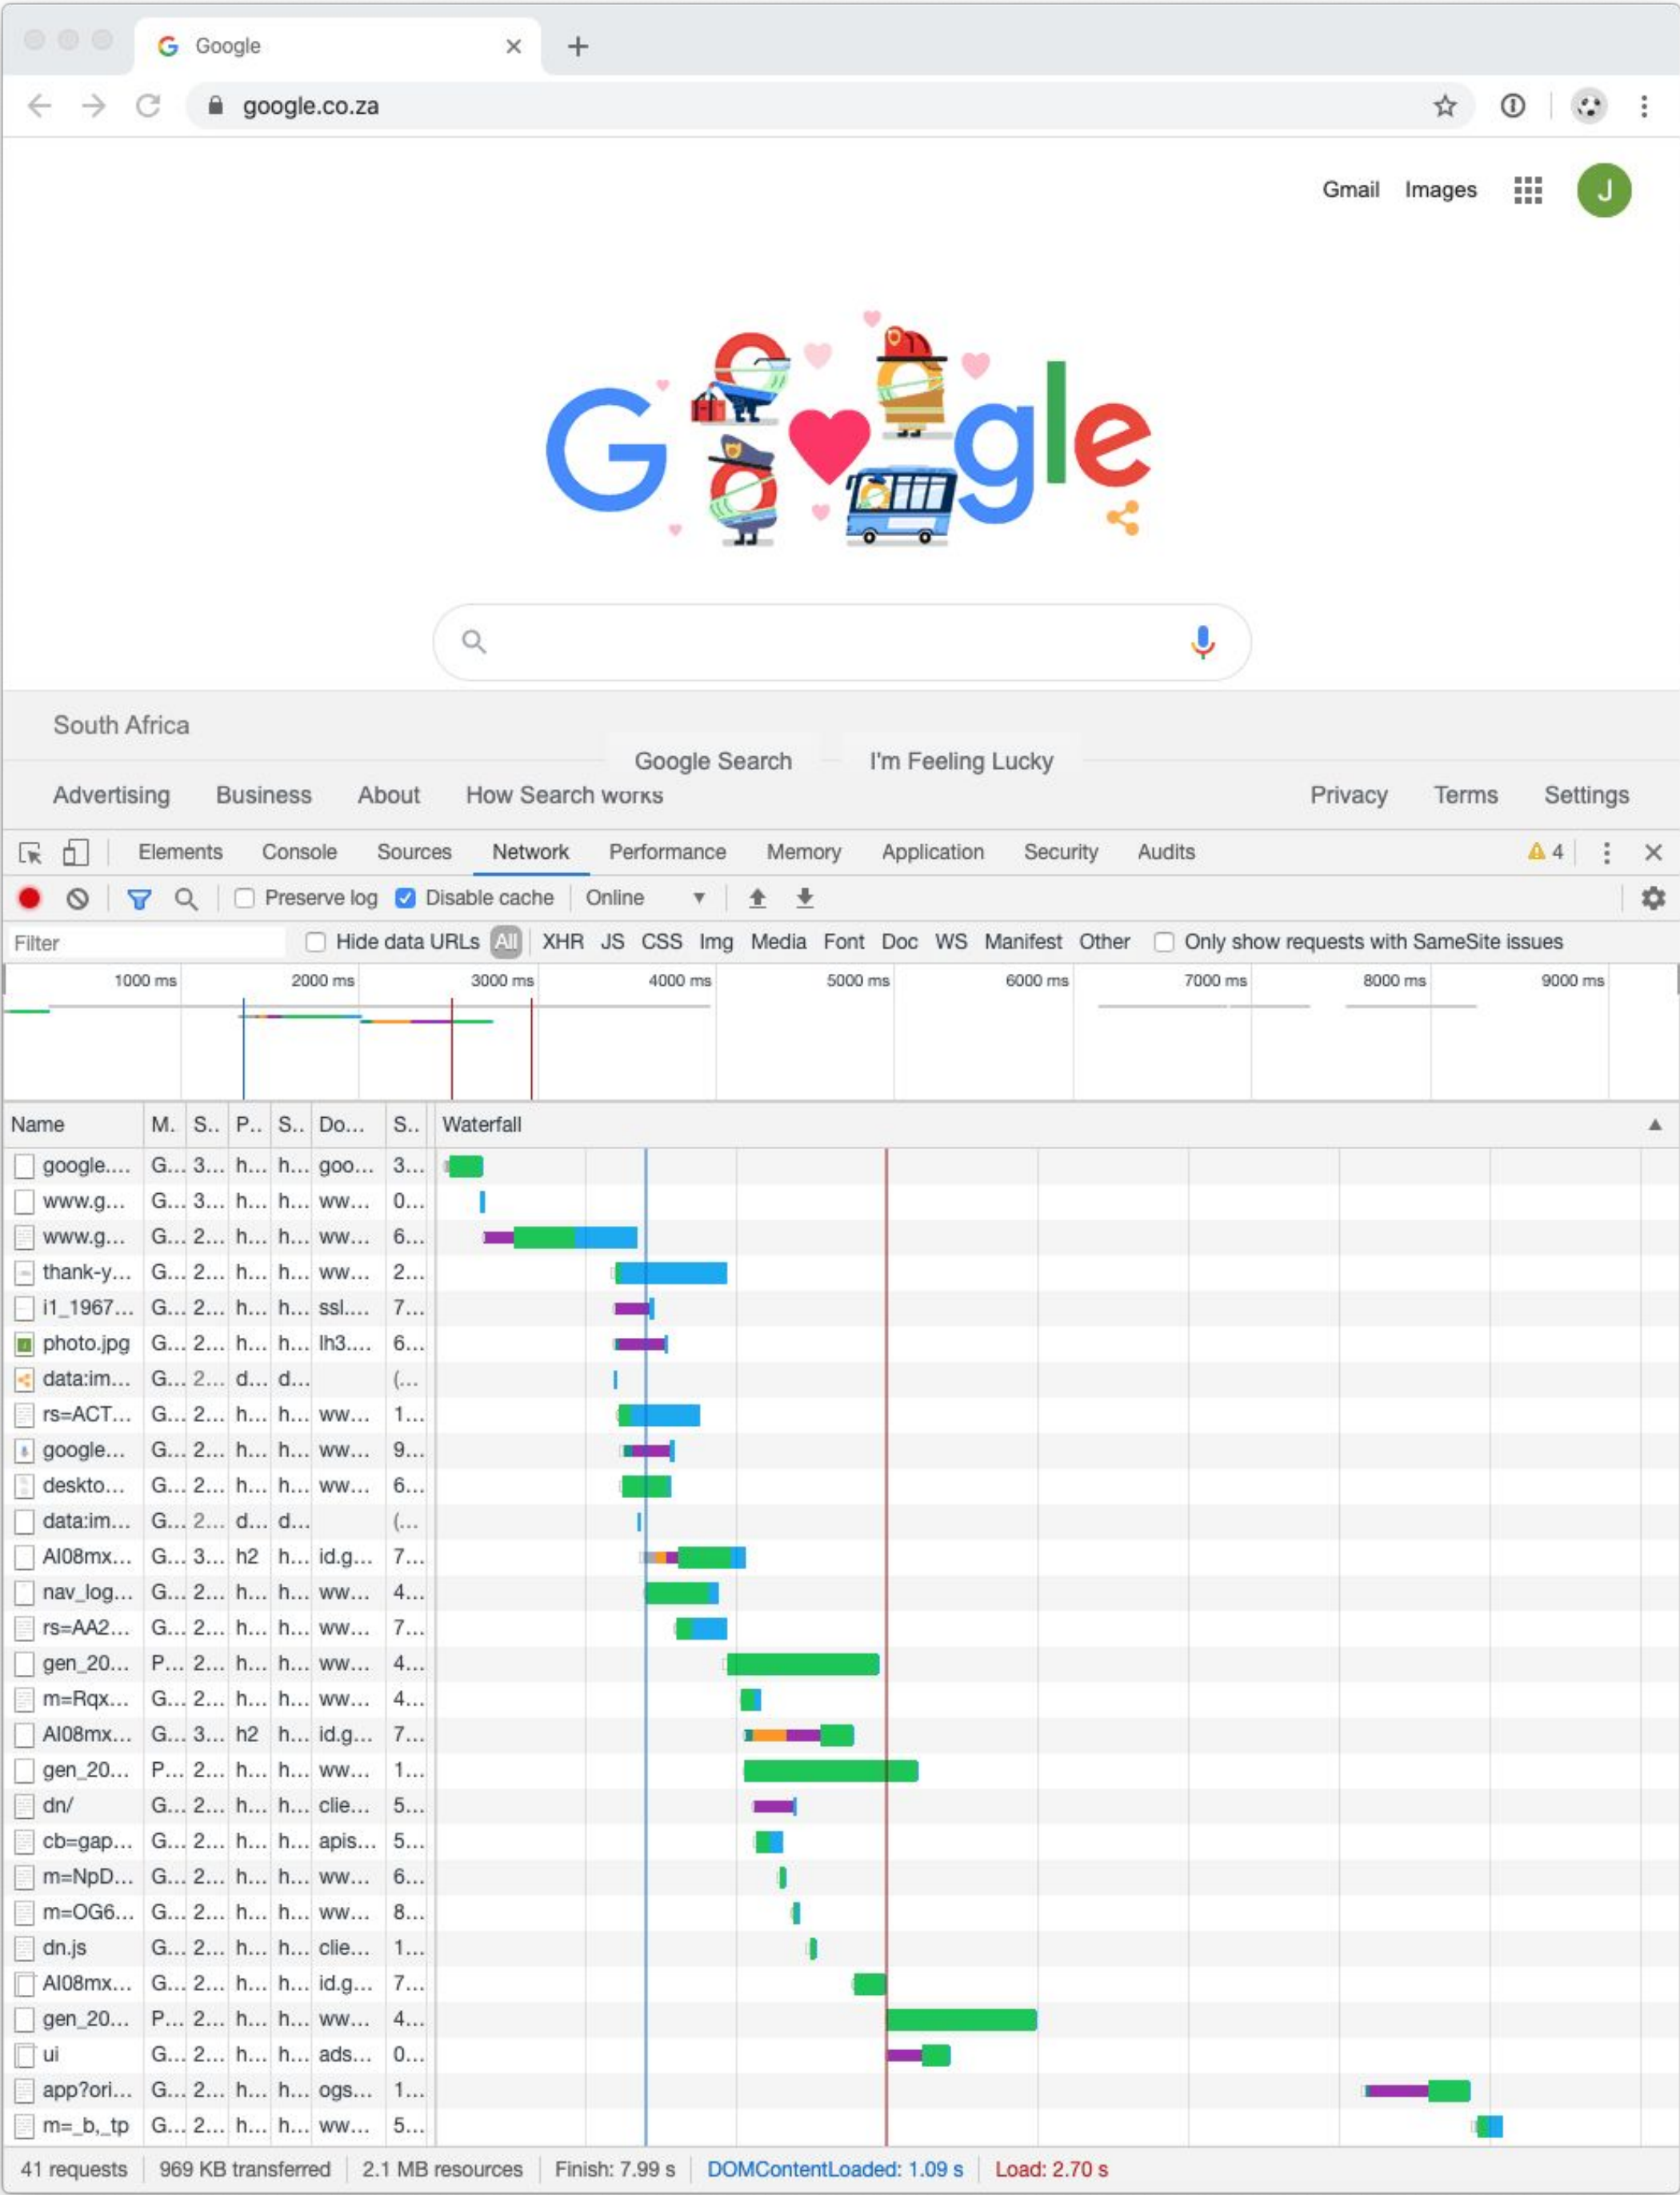
**What if we thought of a Mojaloop transfer as a website?**

**How would we optimize it?**

# The Good...



# The Bad...





# The Mojaloop...



18 Events / 13 ms Sum / 134 ms Total

18 SQL queries in 13ms / 134 ms total transfer time

# Catching SQL Queries

- Instrumented Knex MySQL binding for combined handlers

1 transfer = 11 inserts + 2 updates + 6 selects  
= 18 SQL queries

# Fsync

Tool	Rate	Latency	Notes
SAMSUNG MZ7LN512	160/s	6.3ms	Consumer grade SATA
Crucial_CT480M500SSD1	108/s	9.3ms	Consumer grade SATA
Intel 520	2031/s	0.49ms	Consumer grade SATA
SAMSUNG MZVPV512HDGL	104/s	9.6ms	Consumer grade NVMe
Samsung SSD 960 PRO	267/s	3.8ms	High-end consumer grade NVMe
Intel 750	2038/s	0.49ms	High-end consumer grade NVMe
Intel PC-3700	7380/s	0.14ms	High-end enterprise-grade NVMe

Source: <https://www.percona.com/blog/2018/02/08/fsync-performance-storage-devices>

# ilp-packet

```
insert into `transfer` (`amount`, `currencyId`, `expirationDate`,  
`ilpCondition`, `transferId`) values (?, ?, ?, ?, ?)
```

```
insert into `ilpPacket` (`transferId`, `value`) values (?, ?)
```



# Duplicate Checks

Before:

```
select * from `transferDuplicateCheck` where `transferId` = ?
```

After:

```
insert into `transferDuplicateCheck` (`hash`, `transferId`) values (?, ?)
```

Can we skip both these queries 99% of the time and still detect duplicates safely?

# Duplicate Checks

Before:

```
select * from `transferDuplicateCheck` where `transferId` = ?
```

After:

```
insert into `transferDuplicateCheck` (`hash`, `transferId`) values (?, ?)
```

Can we skip both these queries 99% of the time and still detect duplicates safely?

Yes we can: Bloom filter

# What else do you notice?



18 Events / 13 ms Sum / 134 ms Total

18 SQL queries in 13ms / 134 ms total transfer time



# The Big Gap



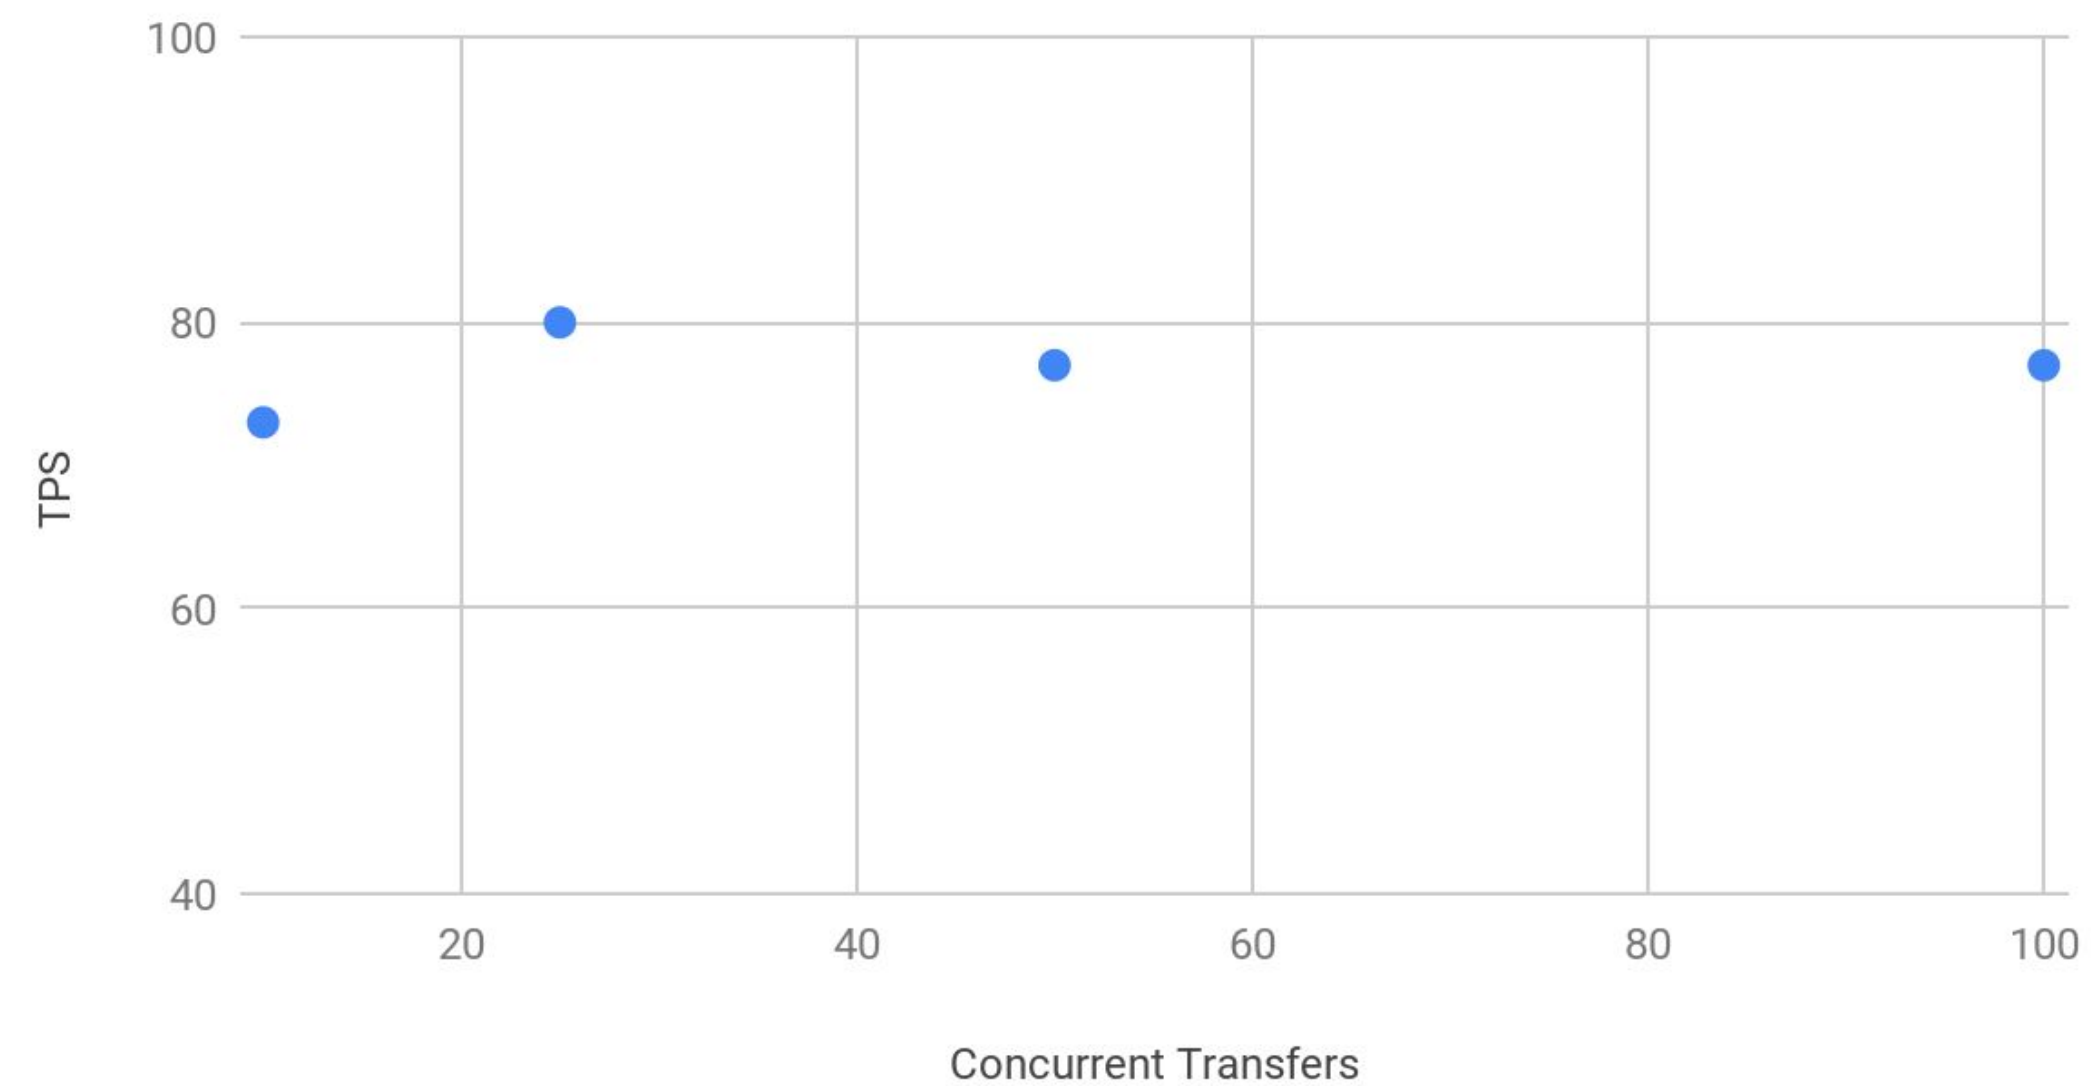
18 Events / 13 ms Sum / 134 ms Total

18 SQL queries in 13ms / 134 ms total transfer time

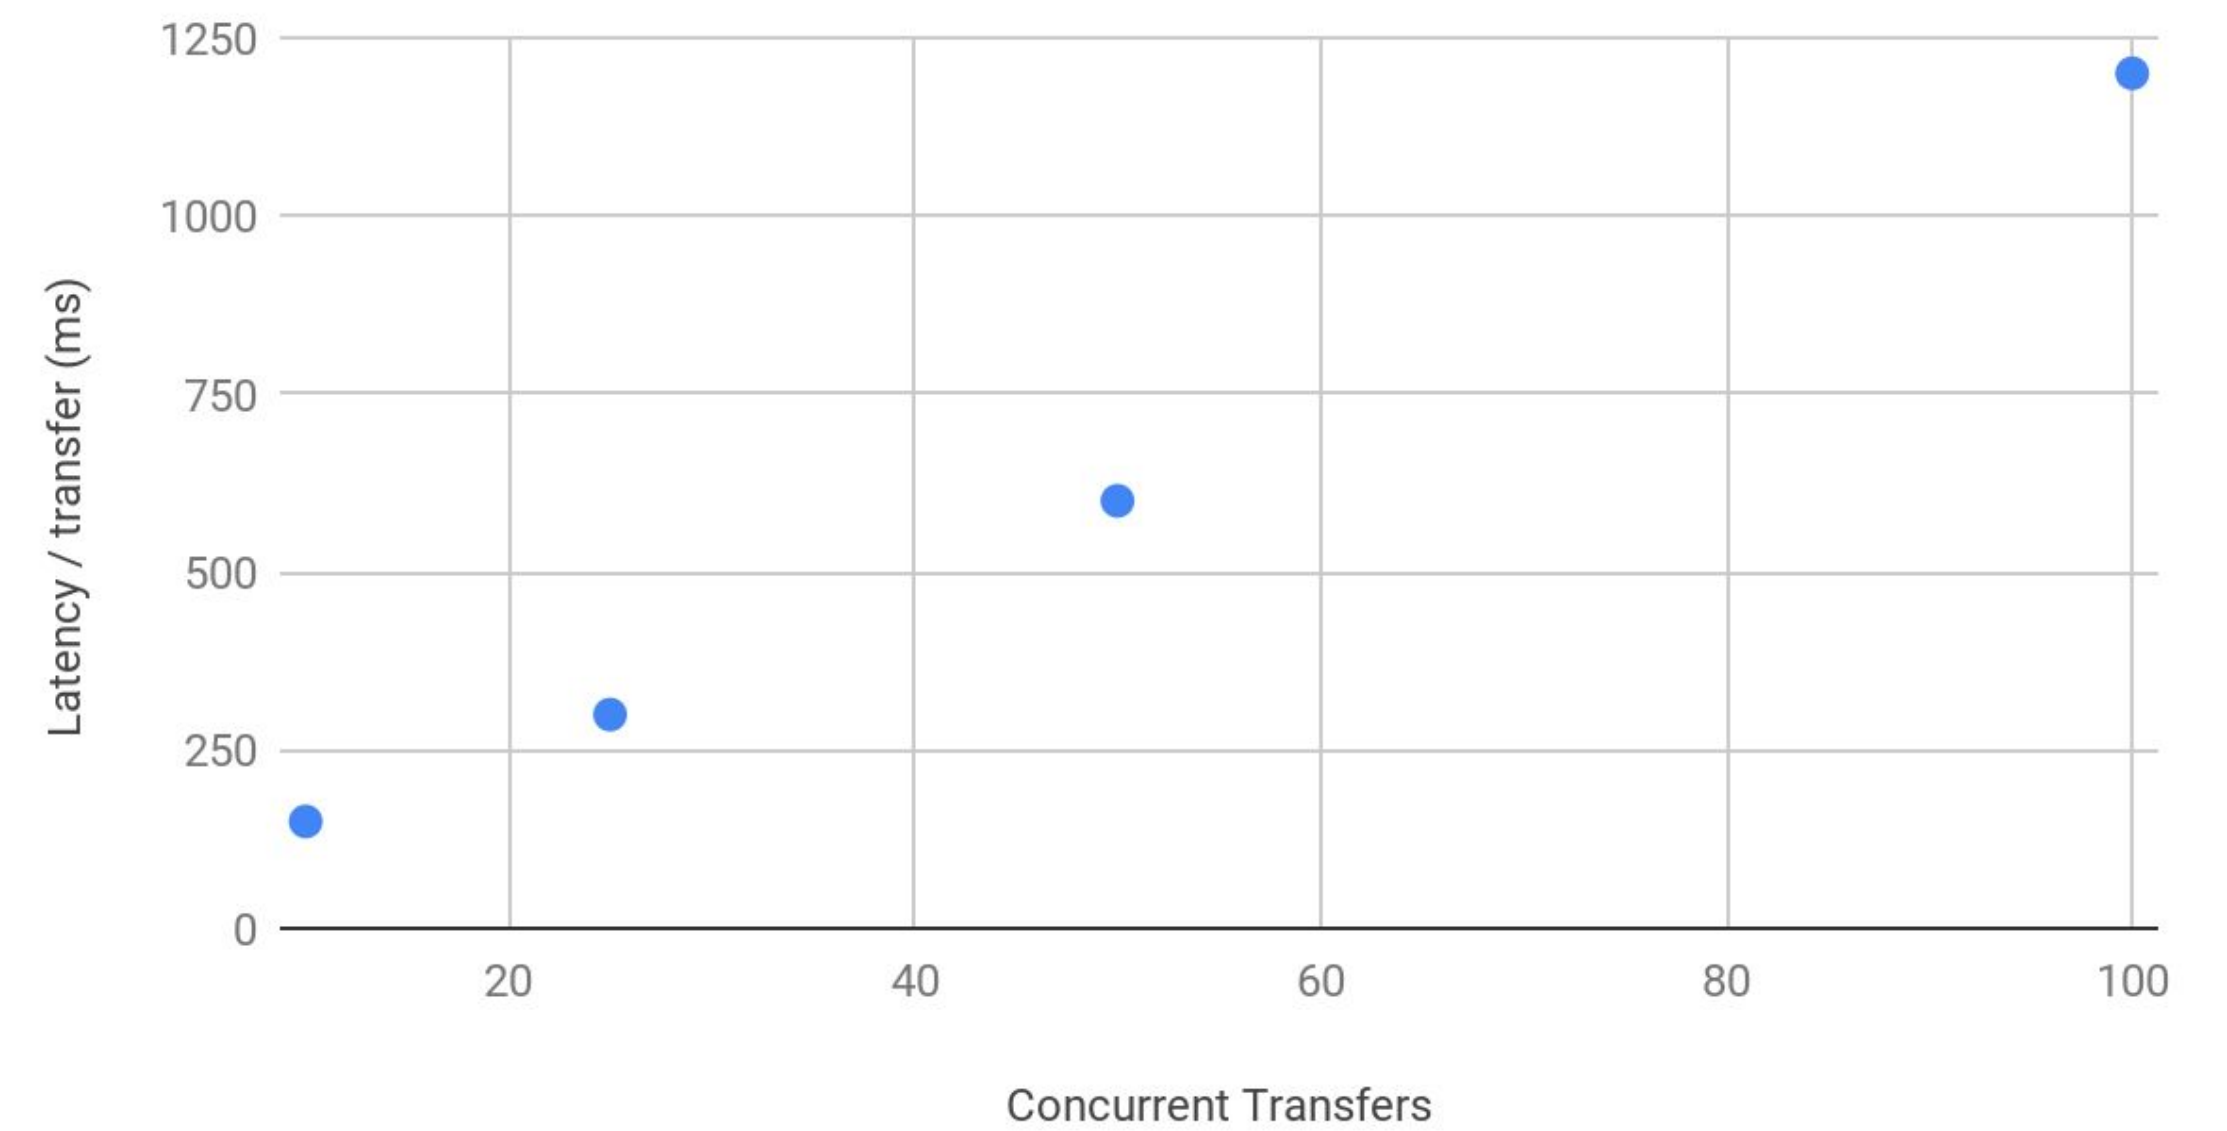


# Concurrency

Throughput vs Concurrent Transfers



Latency / transfer vs. Concurrent Transfers



# 100 Concurrent Transfers

meta: transfer: 275ad688-3f1a-4fa6-96f8-ea1688ed5ea1		1233 ms
meta: prepare notification: 275ad688-3f1a-4fa6-96f8-ea1688ed5ea1		559 ms
meta: prepare/position handler	12 ms	
insert into `transferDuplicateCheck` (`hash`, `transferId`) values (?, ?)	1 ms	
insert into `ilpPacket` (`transferId`, `value`) values (?, ?)	2 ms	
insert into `transferParticipant` (`amount`, `ledgerEntryTypeId`, `particip...	2 ms	
insert into `transferParticipant` (`amount`, `ledgerEntryTypeId`, `particip...	1 ms	
insert into `transfer` (`amount`, `currencyId`, `expirationDate`, `ilpCondi...	1 ms	
insert into `transferStateChange` (`createdDate`, `reason`, `transferId`, ...	1 ms	
INSERT INTO participantPositionChange (participantPositionId, transfe...	1 ms	
UPDATE participantPosition SET value = (value + 100), changedDate ...	0 ms	
meta: fulfil notification: 275ad688-3f1a-4fa6-96f8-ea1688ed5ea1		674 ms
meta: fulfil/position handler	15 ms	
select `transfer`.*, `transfer`.`currencyId` as `currency`, `pc1`.`participa...	2 ms	
select * from `transferExtension` where `transferId` = ? and `isFulfilmen...	1 ms	
insert into `transferFulfilmentDuplicateCheck` (`hash`, `transferId`) valu...	1 ms	
select `transferParticipant`.*, `tsc`.`transferStateId`, `tsc`.`reason` from ...	1 ms	
select `settlementWindow`.`settlementWindowId`, `swsc`.`settlementW...	1 ms	
insert into `transferFulfilment` (`completedDate`, `createdDate`, `ilpFulfi...	1 ms	
insert into `transferStateChange` (`createdDate`, `transferId`, `transfer...	0 ms	
select * from `transferStateChange` where `transferId` = ? order by `tra...	0 ms	
INSERT INTO participantPositionChange (participantPositionId, transfe...	1 ms	
UPDATE participantPosition SET value = (value + -100), changedDate ...	0 ms	

18 Events / 13 ms Sum / 1233 ms Total

18 SQL queries in 13ms / 1233 ms total transfer time

# Next Steps

- Instrument the notification handler
- Instrument kafka calls to produce to and consume off the queue
- Investigate the effect of pulling multiple messages off the queue at the same time, by the handlers

# DNS Lookups

## Bad News

<https://github.com/nodejs/node/issues/8436>

- Node's dns.lookup runs in the threadpool
- Default UV\_THREADPOOL\_SIZE = 4

## Good News

- Monkey-patched DNS lookups
- None were in the critical path



Help us solve the big gap problem!

# Environment

Node type: n1-standard-4 4vCPUs 15GiB RAM

Component	Pod Scale	Dedicated Nodes	label
Kafka + ZooKeeper		3	broker
MySQL	1	1	data
ML-API-Adapter	1	1	ml_api
Prepare handler	1	1	ml_cl_prepare
Position handler	1	1	ml_cl_position
Fulfill handler	1	1	ml_cl_fulfil
Notification handler	1	1	ml_notify
Load generator	1	1	load
monitoring		1	monitor